

# Index the Data Lake storage with Hyperspace

---

When loading data from Azure Data Lake Gen 2, searching in the data is one of the most resource consuming operations. [Hyperspace](#) introduces the ability for Apache Spark users to create indexes on their datasets, such as CSV, JSON, and Parquet, and use them for potential query and workload acceleration.

Hyperspace lets you create indexes on records scanned from persisted data files. After they're successfully created, an entry that corresponds to the index is added to the Hyperspace's metadata. This metadata is later used by Apache Spark's optimizer during query processing to find and use proper indexes. If the underlying data changes, you can refresh an existing index to capture that.

Also, Hyperspace allows users to compare their original plan versus the updated index-dependent plan before running their query.

Add a new cell to your notebook with the following code (remember to replace `<unique_suffix>` with the value you specified during the Synapse Analytics workspace deployment):

```
from hyperspace import *
from com.microsoft.hyperspace import *
from com.microsoft.hyperspace.index import *

# Disable BroadcastHashJoin, so Spark will use standard SortMergeJoin. Currently,
# Hyperspace indexes utilize SortMergeJoin to speed up query.
spark.conf.set("spark.sql.autoBroadcastJoinThreshold", -1)

dfSales = spark.read.parquet("abfss://wwi-
02@asagadatalake<unique_suffix>.dfs.core.windows.net/sale-
small/Year=2019/Quarter=Q4/Month=12/*/*.parquet")
dfSales.show(10)

dfCustomers = spark.read.load("abfss://wwi-
02@asagadatalake<unique_suffix>.dfs.core.windows.net/data-generators/generator-
customer-clean.csv", format="csv", header=True)
dfCustomers.show(10)

# Create an instance of Hyperspace
hyperspace = Hyperspace(spark)
```

Run the new cell. It will load the two DataFrames with data from the data lake and initialize Hyperspace.

```

1 from hyperspace import *
2 from com.microsoft.hyperspace import *
3 from com.microsoft.hyperspace.index import *
4
5 # Disable BroadcastHashJoin, so Spark will use standard SortMergeJoin. Currently, Hyperspace indexes utilize SortMergeJoin to speed up query.
6 spark.conf.set("spark.sql.autoBroadcastJoinThreshold", -1)
7
8 dfSales = spark.read.parquet("abfss://wwi-02@asadatalake01.dfs.core.windows.net/sale-small/Year=2019/Quarter=Q4/Month=12/*/*.parquet")
9 dfSales.show(10)
10
11 dfCustomers = spark.read.load("abfss://wwi-02@asadatalake01.dfs.core.windows.net/data-generators/generator-customer-clean.csv", format="csv", header=True)
12 dfCustomers.show(10)
13
14 # Create an instance of Hyperspace
15 hyperspace = Hyperspace(spark)

```

Command executed in 23s 590ms by ciprian on 11-30-2020 22:43:56.200 +02:00

Job execution Succeeded Spark 2 executors 8 cores [View in monitoring](#) [Open Spark UI](#)

TransactionId	CustomerId	ProductId	Quantity	Price	TotalAmount	TransactionDate	ProfitAmount	Hour	Minute	StoreId
7bcfb8f0-d3ec-4fa...	2	828	2	213.28000000000000	166.56000000000000	20191204	12.66000000000000	6	2	92
7bcfb8f0-d3ec-4fa...	2	1370	2	26.66000000000000	53.32000000000000	20191204	16.58000000000000	6	2	92
7bcfb8f0-d3ec-4fa...	2	3071	1	139.22000000000000	39.22000000000000	20191204	12.38000000000000	6	2	92
7bcfb8f0-d3ec-4fa...	2	3460	4	30.84000000000000	123.36000000000000	20191204	42.72000000000000	6	2	92
7bcfb8f0-d3ec-4fa...	2	2975	1	133.24000000000000	33.24000000000000	20191204	11.19000000000000	6	2	92
7bcfb8f0-d3ec-4fa...	2	4889	2	32.02000000000000	64.04000000000000	20191204	20.20000000000000	6	2	92
7bcfb8f0-d3ec-4fa...	2	2333	1	119.22000000000000	19.22000000000000	20191204	5.54000000000000	6	2	92
7bcfb8f0-d3ec-4fa...	2	105	1	134.36000000000000	34.36000000000000	20191204	10.51000000000000	6	2	92
7bcfb8f0-d3ec-4fa...	2	6	1	128.31000000000000	28.31000000000000	20191204	8.77000000000000	6	2	92
7bcfb8f0-d3ec-4fa...	2	144	3	30.44000000000000	91.44000000000000	20191204	24.79000000000000	6	2	92

only showing top 10 rows

CustomerId	FirstName	MiddleInitial	LastName	FullName	Gender	Age	BirthDate	Address_PostalCode	Address_Street	Address_City	Address_Country	Mobile	Email
1	Conrad	N	Zemlak	Conrad N. Zemlak	Male	57	1963-03-23	16219	Dario View	Dockside	Suriname	0925-586-826	Conrad_Zemlak_85@...
2	Darren	H	Rolfson	Darren H. Rolfson	Male	61	1959-05-25	43387	Kathleen Oval	Lake Rileytown	Tanzania	0372-253-946	Darren_Rolfson_20...
3	Annie	J	Pouros	Annie J. Pouros	Female	51	1969-04-23	92666	Maegan Hill	Lake Eltchester	Virgin Islands Br...	0834-564-182	Annie_Pouros_48@y...
4	Forrest	null	O\Kon	Forrest O\Kon	Male	54	1966-01-31	29524	Harry Spring	New Gregoriaport	Netherlands	0761-327-187	Forrest_O\Kon_83...
5	Leroy	null	Lemke	Leroy Lemke	Male	59	1961-03-01	36967	Layla Ford	East Davontebury	Svalbard & Jan Ma...	0264-161-776	Leroy_Lemke_29@ou...
6	Luz	R	Borer	Luz R. Borer	Female	53	1967-04-28	51050	Marjorie Overpass	Carleeshire	Burkina Faso	0642-895-538	Luz_Borer_77@outl...
7	Ira	M	Koelpin	Ira M. Koelpin	Male	52	1968-04-22	69710	Kendall Street	New Lola	Suriname	0114-516-150	Ira_Koelpin_60@gm...
8	Salvatore	null	Schaden	Salvatore Schaden	Male	54	1965-12-30	45947	Juwan Crest	Thompsonhaven	Liechtenstein	0782-905-529	Salvatore_Schaden...
9	Robyn	K	Batz	Robyn K. Batz	Female	59	1960-08-17	67220	Derrick Flat	Lake Joeymouth	Denmark	0729-374-410	Robyn_Batz_10@ma...
10	Brandi	D	Kertzmann	Brandi D. Kertzmann	Female	43	1976-12-02	44784	Madisen Track	East Hyriamhaven	Moldova	0739-782-555	Brandi_Kertzmann...

only showing top 10 rows

Add another new cell to your notebook with the following code:

```

#create indexes: each one contains a name, a set of indexed columns and a set of included columns
indexConfigSales = IndexConfig("indexSALES", ["CustomerId"], ["TotalAmount"])
indexConfigCustomers = IndexConfig("indexCUSTOMERS", ["CustomerId"], ["FullName"])

hyperspace.createIndex(dfSales, indexConfigSales) # only create index once
hyperspace.createIndex(dfCustomers, indexConfigCustomers) # only create index once
hyperspace.indexes().show()

```

Run the new cell. It will create two indexes and display their structure.

```

1 #create indexes: each one contains a name, a set of indexed columns and a set of included columns
2 indexConfigSales = IndexConfig("indexSALES", ["CustomerId"], ["TotalAmount"])
3 indexConfigCustomers = IndexConfig("indexCUSTOMERS", ["CustomerId"], ["FullName"])
4
5 hyperspace.createIndex(dfSales, indexConfigSales) # only create index once
6 hyperspace.createIndex(dfCustomers, indexConfigCustomers) # only create index once
7 hyperspace.indexes().show()

```

Command executed in 49s 512ms by ciprian on 11-30-2020 22:49:09.292 +02:00

Job execution Succeeded Spark 2 executors 8 cores [View in monitoring](#)

name	indexedColumns	includedColumns	numBuckets	schema	indexLocation	state
indexCUSTOMERS	[CustomerId]	[FullName]	200	{ "type": "struct", ... }	abfss://workspace...	ACTIVE
indexSALES	[CustomerId]	[TotalAmount]	200	{ "type": "struct", ... }	abfss://workspace...	ACTIVE

Add another new cell to your notebook with the following code:

```

df1 = dfSales.filter("CustomerId = 203").select("TotalAmount")
df1.show()
df1.explain(True)

```

Run the new cell. The output will show that the physical execution plan is not taking into account any of the indexes (performs a file scan on the original data file).

```

1 df1 = dfSales.filter("""CustomerId = 203""").select("""TotalAmount""")
2 df1.show()
3 df1.explain(True)

```

Command executed in 15s 308ms by ciprian on 11-30-2020 23:50:33.836 +02:00

Job execution Succeeded Spark 2 executors 8 cores [View in monitoring](#) [Open Spark UI](#)

```

+-----+
|      TotalAmount      |
+-----+
|28.330000000000000000...|
|110.9100000000000000...|
|95.7200000000000000...|
|103.8400000000000000...|
|133.3200000000000000...|
|92.2800000000000000...|
|31.2800000000000000...|
+-----+

== Parsed Logical Plan ==
'Project [unresolvedalias('TotalAmount, None)]
+- Filter (CustomerId#28 = 203)
   +- Relation[TransactionId#27,CustomerId#28,ProductId#29,Quantity#30,Price#31,TotalAmount#32,TransactionDate#33,ProfitAmount#34,Hour#35,Minute#36,StoreId#37] parquet

== Analyzed Logical Plan ==
TotalAmount: decimal(38,18)
Project [TotalAmount#32]
+- Filter (CustomerId#28 = 203)
   +- Relation[TransactionId#27,CustomerId#28,ProductId#29,Quantity#30,Price#31,TotalAmount#32,TransactionDate#33,ProfitAmount#34,Hour#35,Minute#36,StoreId#37] parquet

== Optimized Logical Plan ==
Project [TotalAmount#32]
+- Filter (isNotNull(CustomerId#28) && (CustomerId#28 = 203))
   +- Relation[TransactionId#27,CustomerId#28,ProductId#29,Quantity#30,Price#31,TotalAmount#32,TransactionDate#33,ProfitAmount#34,Hour#35,Minute#36,StoreId#37] parquet

== Physical Plan ==
*(1) Project [TotalAmount#32]
+- *(1) Filter (isNotNull(CustomerId#28) && (CustomerId#28 = 203))
   +- *(1) FileScan parquet [CustomerId#28,TotalAmount#32] Batched: true, Format: Parquet, Location: InMemoryFileIndex[abfss://ww1-02@asadatalake01.dfs.core.windows.net/sale-small/Year=2019/Quarter=...], PartitionFilters: [], PushedFilters: [IsNotNull(CustomerId), EqualTo(CustomerId,203)], ReadSchema: struct<CustomerId:int,TotalAmount:decimal(38,18)>

```

Now add another new cell to your notebook with the following code (notice the extra line at the beginning used to enable Hyperspace optimization in the Spark engine):

```

# Enable Hyperspace - Hyperspace optimization rules become visible to the Spark
optimizer and exploit existing Hyperspace indexes to optimize user queries
Hyperspace.enable(spark)
df1 = dfSales.filter("""CustomerId = 203""").select("""TotalAmount""")
df1.show()
df1.explain(True)

```

Run the new cell. The output will show that the physical execution plan is now using the index instead of the original data file.

```

1 # Enable Hyperspace - Hyperspace optimization rules become visible to the Spark
2 Hyperspace.enable(spark)
3 df1 = dfSales.filter("""CustomerId = 203""").select("""TotalAmount""")
4 df1.show()
5 df1.explain(True)

```

Command executed in 13s 417ms by ciprian on 11-30-2020 23:54:56.589 +02:00

Job execution Succeeded Spark 2 executors 8 cores [View in monitoring](#) [Open Spark UI](#)

```

+-----+
|      TotalAmount      |
+-----+
|28.330000000000000000...|
|110.9100000000000000...|
|95.7200000000000000...|
|103.8400000000000000...|
|133.3200000000000000...|
|92.2800000000000000...|
|31.2800000000000000...|
+-----+

== Parsed Logical Plan ==
'Project [unresolvedalias('TotalAmount, None)]
+- Filter (CustomerId#28 = 203)
   +- Relation[TransactionId#27,CustomerId#28,ProductId#29,Quantity#30,Price#31,TotalAmount#32,TransactionDate#33,ProfitAmount#34,Hour#35,Minute#36,StoreId#37] parquet

== Analyzed Logical Plan ==
TotalAmount: decimal(38,18)
Project [TotalAmount#32]
+- Filter (CustomerId#28 = 203)
   +- Relation[TransactionId#27,CustomerId#28,ProductId#29,Quantity#30,Price#31,TotalAmount#32,TransactionDate#33,ProfitAmount#34,Hour#35,Minute#36,StoreId#37] parquet

== Optimized Logical Plan ==
Project [TotalAmount#32]
+- Filter (isNotNull(CustomerId#28) && (CustomerId#28 = 203))
   +- Relation[CustomerId#28,TotalAmount#32] parquet

== Physical Plan ==
*(1) Project [TotalAmount#32]
+- *(1) Filter (isNotNull(CustomerId#28) && (CustomerId#28 = 203))
   +- *(1) FileScan parquet [CustomerId#28,TotalAmount#32] Batched: true, Format: Parquet, Location: InMemoryFileIndex[abfss://workspace@asadatalake01.dfs.core.windows.net/synapse/workspaces/asaga...], PartitionFilters: [], PushedFilters: [IsNotNull(CustomerId), EqualTo(CustomerId,203)], ReadSchema: struct<CustomerId:int,TotalAmount:decimal(38,18)>

```

Hyperspace provides an Explain API that allows you to compare the execution plans without indexes vs. with indexes. Add a new cell with the following code:

```
df1 = dfSales.filter("""CustomerId = 203""").select("""TotalAmount""")

spark.conf.set("spark.hyperspace.explain.displayMode", "html")
hyperspace.explain(df1, True, displayHTML)
```

Run the new cell. The output shows a comparison **Plan with indexes** vs. **Plan without indexes**. Observe how, in the first case the index file is used while in the second case the original data file is used.

```
1 df1 = dfSales.filter("""CustomerId = 203""").select("""TotalAmount""")
2
3 spark.conf.set("spark.hyperspace.explain.displayMode", "html")
4 hyperspace.explain(df1, True, displayHTML)
```

Command executed in 3s 924ms by ciprian on 12-01-2020 00:00:07.548 +02:00

> Job execution Succeeded Spark 2 executors 8 cores [View in monitoring](#) [Open Spark UI](#)

```
=====
Plan with indexes:
=====
Project [TotalAmount#32]
+- Filter (isNotNull(CustomerId#28) && (CustomerId#28 = 203))
   +- FileScan parquet [CustomerId#28,TotalAmount#32] Batched: true, Format: Parquet, Location: InMemoryFileIndex[abfss://workspace@asagadatalake01.dfs.core.windows.net/synapse/workspaces/asaga..., PartitionFi

=====
Plan without indexes:
=====
Project [TotalAmount#32]
+- Filter (isNotNull(CustomerId#28) && (CustomerId#28 = 203))
   +- FileScan parquet [CustomerId#28,TotalAmount#32] Batched: true, Format: Parquet, Location: InMemoryFileIndex[abfss://wwi-02@asadatalake01.dfs.core.windows.net/sale-small/Year=2019/Quarter=..., PartitionFi

=====
Indexes used:
=====
indexSALES:abfss://workspace@asagadatalake01.dfs.core.windows.net/synapse/workspaces/asagaworkspace01/warehouse/indexes/indexSALES/v__=0

=====
Physical operator stats:
=====
+-----+
|Physical Operator|Hyperspace Disabled|Hyperspace Enabled|Difference|
+-----+
|Filter|1|1|0|
|Project|1|1|0|
|Scan parquet|1|1|0|
|WholeStageCodegen|1|1|0|
+-----+
```

Let's investigate now a more complex case, involving a join operation. Add a new cell with the following code:

```
eqJoin = dfSales.join(dfCustomers, dfSales.CustomerId ==
dfCustomers.CustomerId).select(dfSales.TotalAmount, dfCustomers.FullName)

hyperspace.explain(eqJoin, True, displayHTML)
```

Run the new cell. The output shows again a comparison **Plan with indexes** vs. **Plan without indexes**, where indexes are used in the first case and the original data files in the second.

```

1 eqJoin = dfSales.join(dfCustomers, dfSales.CustomerId == dfCustomers.CustomerId).select(dfSales.TotalAmount, dfCustomers.FullName)
2
3 hyperspace.explain(eqJoin, True, displayHTML)

```

Command executed in Scala by opspark on 12-01-2020 00:21:16.515 - 02:00

Job execution Succeeded Spark 2 executors 8 cores [View in monitoring](#) [Open](#)

=====

Plan with Indexes:

```

Project [TotalAmount#32, FullName#97]
+- SortMergeJoin [CustomerId#28], [cast(CustomerId#93 as int)], Inner
   +- *(2) Sort [CustomerId#28 ASC NULLS FIRST], false, 0
      +- Exchange hashpartitioning(CustomerId#28, 200), [id=#1487]
         +- *(1) Project [CustomerId#28, TotalAmount#32]
            +- *(2) Filter [isNotNull(CustomerId#28)]
               +- *(1) FileScan parquet [CustomerId#28,TotalAmount#32] Batched: true, Format: Parquet, Location: InMemoryFileIndex[abfss://workspace@asadatalake01.dfs.core.windows.net/synapse/workspaces/asga..., PartitionFilters: [], PushedFilters: [IsNotNull(CustomerId)], ReadSchema: struct
                  +- *(4) Sort [cast(CustomerId#93 as int) ASC NULLS FIRST], false, 0
                     +- Exchange hashpartitioning(cast(CustomerId#93 as int), 200), [id=#1493]
                        +- *(3) Project [CustomerId#93, FullName#97]
                           +- *(2) Filter [isNotNull(CustomerId#93)]
                              +- *(3) FileScan csv [CustomerId#93,FullName#97] Batched: true, Format: CSV, Location: InMemoryFileIndex[abfss://workspace@asadatalake01.dfs.core.windows.net/synapse/workspaces/asga..., PartitionFilters: [], PushedFilters: [IsNotNull(CustomerId)], ReadSchema: struct

```

=====

Plan without Indexes:

```

Project [TotalAmount#32, FullName#97]
+- SortMergeJoin [CustomerId#28], [cast(CustomerId#93 as int)], Inner
   +- *(2) Sort [CustomerId#28 ASC NULLS FIRST], false, 0
      +- Exchange hashpartitioning(CustomerId#28, 200), [id=#1450]
         +- *(1) Project [CustomerId#28, TotalAmount#32]
            +- *(1) Filter [isNotNull(CustomerId#28)]
               +- *(1) FileScan parquet [CustomerId#28,TotalAmount#32] Batched: true, Format: Parquet, Location: InMemoryFileIndex[abfss://uni-02@asadatalake01.dfs.core.windows.net/sale-small/year=2019/Quarters..., PartitionFilters: [], PushedFilters: [IsNotNull(CustomerId)], ReadSchema: struct
                  +- *(4) Sort [cast(CustomerId#93 as int) ASC NULLS FIRST], false, 0
                     +- Exchange hashpartitioning(cast(CustomerId#93 as int), 200), [id=#1456]
                        +- *(3) Project [CustomerId#93, FullName#97]
                           +- *(3) Filter [isNotNull(CustomerId#93)]
                              +- *(1) FileScan csv [CustomerId#93,FullName#97] Batched: false, Format: CSV, Location: InMemoryFileIndex[abfss://uni-02@asadatalake01.dfs.core.windows.net/data-generators/generator-cus..., PartitionFilters: [], PushedFilters: [IsNotNull(CustomerId)], ReadSchema: struct

```

=====

Indexes used:

```

IndexCUSTOMERS:abfss://workspace@asadatalake01.dfs.core.windows.net/synapse/workspaces/asga/workspace01/warehouse/Indexes/indexCUSTOMERS/v_...0
IndexSALES:abfss://workspace@asadatalake01.dfs.core.windows.net/synapse/workspaces/asga/workspace01/warehouse/Indexes/indexSALES/v_...0

```

=====

Physical operator stats:

```

-----+-----+
|Physical Operator|Hyperspace Disabled|Hyperspace Enabled|Difference|
|-----+-----+-----+-----+
|*Scan csv|1|0|-1|
|*Scan parquet|1|2|1|
|Filter|2|0|-2|
|InputAdapter|4|4|0|
|Project|3|3|0|
|ShuffleExchange|2|2|0|
|Sort|2|2|0|
|SortMergeJoin|1|1|0|
|WholeStageCodeGen|5|5|0|
|-----+-----+-----+-----+

```

In case you want to deactivate Hyperspace and cleanup the indexes, you can run the following code:

```
# Disable Hyperspace - Hyperspace rules no longer apply during query optimization.
Disabling Hyperspace has no impact on created indexes because they remain intact
Hyperspace.disable(spark)
```

```
hyperspace.deleteIndex("indexSALES")
hyperspace.vacuumIndex("indexSALES")
hyperspace.deleteIndex("indexCUSTOMERS")
hyperspace.vacuumIndex("indexCUSTOMERS")
```